

🔍 California Housing Price Prediction – Full Stack ML Project Summary

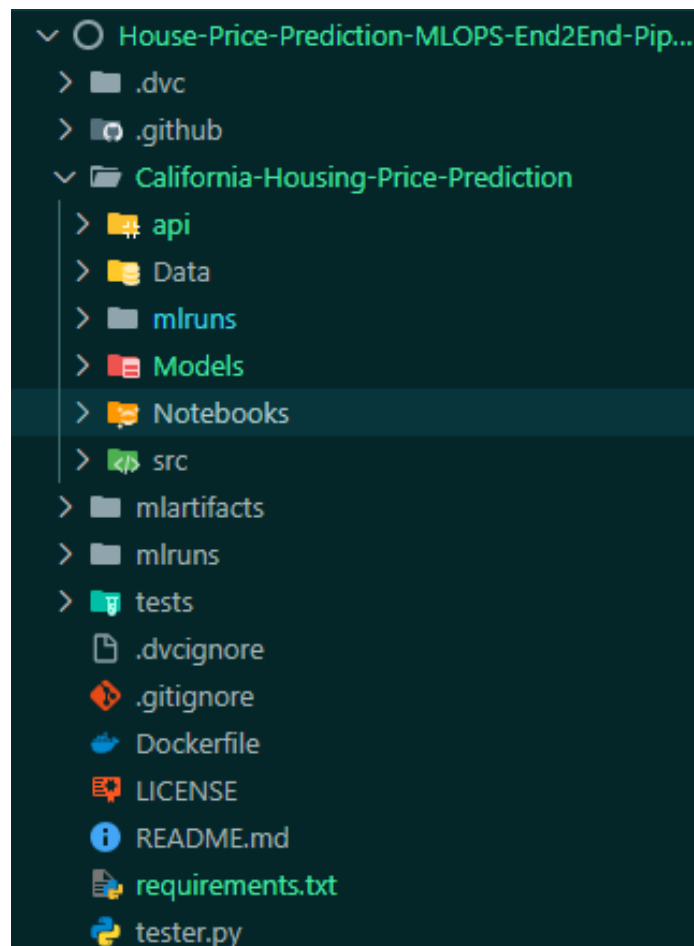
■ Assignment Structure Overview

Project Details:-

GitHub Repo Link: - <https://github.com/Atharv-Chaudhari-Bits/House-Price-Prediction-MLOPS-End2End-Pipeline>

Docker Hub Link: - <https://hub.docker.com/r/atharvchaudharibits/california-housing-app>

Directory Structure –



Built and deployed an end-to-end California Housing Price Prediction pipeline using MLOps best practices. Implemented model training, tracking, and versioning with Git, DVC, and MLflow. Packaged the model as a REST API using FastAPI and containerized it with Docker. Automated CI/CD via GitHub Actions and set up basic logging with optional monitoring support.

✦ ML Models Implemented

We trained and compared several regression models using scikit-learn:

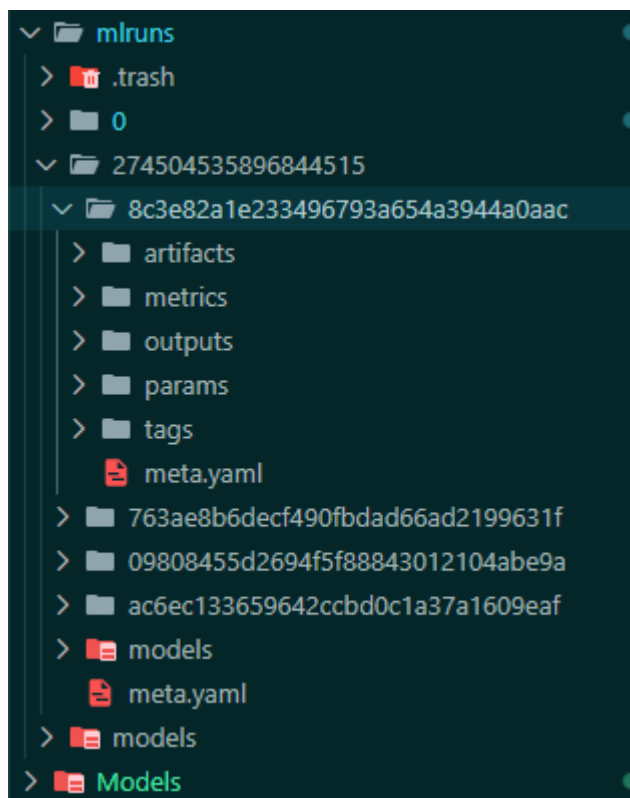
- `LinearRegression`
- `DecisionTreeRegressor`
- `RandomForestRegressor`
- `GradientBoostingRegressor`

Each model was evaluated and logged using **MLflow**, capturing:

- Hyperparameters
- Metrics (R^2 , RMSE, MAE)
- Model artifacts
- Training metadata

MLflow created separate directories under `mlruns/<experiment-id>/<run-id>/`, which include:

- `params/`: model hyperparameters
- `metrics/`: performance metrics
- `artifacts/`: saved models
- `tags/`: metadata like model name or author



🐳 Docker Setup (Model Packaging + API)

The project is fully containerized using **Docker**, enabling smooth deployment across environments.

Key files used:

- `Dockerfile`: Defines the build, installing FastAPI, MLflow, scikit-learn, etc.
- `.dockerignore`: Ignores unnecessary files like `__pycache__`, notebooks, and MLflow trash.
- Docker builds the image with API code and launches a FastAPI server.

Main functionalities containerized:

1. Load the best MLflow model from registry
2. Serve predictions via REST (JSON) and HTML form (`/form`)
3. Logging for inference time, status, and model used

🌐 FastAPI Features

FastAPI is used to expose the model through:

- `/predict`: accepts JSON inputs and returns prediction
- `/form`: accepts user inputs via HTML form and returns output
- `/retrain`: accepts new CSV uploads for dynamic model retraining

Pydantic schemas are used for request validation.

All logs (including prediction attempts) are enhanced with **emojis/symbols** for better visibility in both console and file logs.

Logs are saved to `app_log.log` file

🔄 CI/CD with GitHub Actions

In `.github/workflows/`, I've set up a GitHub Actions pipeline that:

- Installs dependencies
- Runs test suite

📁 Data Versioning with DVC

DVC tracks both raw and processed data (`Data/raw`, `Data/processed`) and integrates with Git:

- Ensures reproducibility by locking datasets to specific model versions
- DVC files (`.dvc/`) and `dvc.yaml` track preprocessing and training steps
- Supports pushing large data to remote (e.g., S3, GDrive)

📦 MLflow Tracking and Model Management

MLflow tracks:

- Training experiments
- Model performance
- Parameters and artifacts
- Registered models under unique IDs

✓ Logging and Monitoring

Logging is implemented for:

- API requests
- Inference status
- Retraining results

Logs include timestamps, status symbols (✓, ✗), and are stored in both console and file format.

⚙️ Retraining Pipeline via API

The `/retrain` endpoint supports:

- Uploading a new CSV via HTML form
- Reading and preprocessing new data
- Retraining the best model
- Logging it to MLflow
- Saving the new model to registry

🧪 Testing

Setup for `tests/` directory includes:

- Unit tests for preprocessing and inference
- API integration tests using `pytest`
- Retraining test cases for model updates

```

=> [internal] load build definition from Dockerfile                                0.2s
=> => transferring dockerfile: 536B                                              0.1s
=> [internal] load metadata for docker.io/library/python:3.12-slim              5.7s
=> [auth] library/python:pull token for registry-1.docker.io                   0.0s
=> [internal] load .dockerignore                                                 0.1s
=> => transferring context: 28                                                    0.0s
=> [1/5] FROM docker.io/library/python:3.12-slim@sha256:9c1d9ed7593f2552a4ea47362ec8d2ddf5923458a53d0c8e30edf8b398c94a31 37.5s
=> => resolve docker.io/library/python:3.12-slim@sha256:9c1d9ed7593f2552a4ea47362ec8d2ddf5923458a53d0c8e30edf8b398c94a31 0.1s
=> => sha256:f5cc5422ebcbbf01f9cd227d36de9dd7e133e1fc6d852f3b0c65260ab58f99f3 250B / 250B 0.0s
=> => sha256:4c665aba06d1c52829be04ca62e1030e27b8a3aa0f922666cbe74d24234ff227 3.51MB / 3.51MB 12.0s
=> => sha256:e3586b415667d044c3e5c7c91023d29d7db667b73a8082068a1b7f36c1962c34 13.66MB / 13.66MB 32.8s
=> => sha256:59e22667830bf04fb35e15ed9c70023e9d121719bb87f0db7f3159ee7c7e0b8d 28.23MB / 28.23MB 30.9s
=> => extracting sha256:59e22667830bf04fb35e15ed9c70023e9d121719bb87f0db7f3159ee7c7e0b8d 3.6s
=> => extracting sha256:4c665aba06d1c52829be04ca62e1030e27b8a3aa0f922666cbe74d24234ff227 0.3s
=> => extracting sha256:e3586b415667d044c3e5c7c91023d29d7db667b73a8082068a1b7f36c1962c34 1.3s
=> => extracting sha256:f5cc5422ebcbbf01f9cd227d36de9dd7e133e1fc6d852f3b0c65260ab58f99f3 0.0s
=> [internal] load build context                                                 5.4s
=> => transferring context: 26.02MB                                              5.3s
=> [2/5] WORKDIR /app                                                           1.0s
=> [3/5] RUN apt-get update && apt-get install -y gcc && rm -rf /var/lib/apt/lists/* 115.4s
=> [4/5] COPY . .                                                                0.6s
=> [5/5] RUN pip install --no-cache-dir --upgrade pip && pip install --no-cache-dir -r requirements.txt 1186.7s
=> exporting to image                                                            562.2s
=> => exporting layers                                                            348.4s
=> => exporting manifest sha256:c801cc2f578d60e73b46c4503aae5d3e7bfc762d55b45966de70d4aaf6cf204 0.1s
=> => exporting config sha256:b32876d308b44640f35db7aee3e9705a481f8e52886474ab0d499506a14fe2c0 0.1s
=> => exporting attestation manifest sha256:050925223b611eea4687990be2fb6a64bfd606c00eecd7f1fd64e713ab183b52 0.1s
=> => exporting manifest list sha256:5233fc527714e196fa5b0b1626eabffe35d753b9f62c54af952a90a05de38765 0.0s
=> => naming to docker.io/library/california-housing-app:latest                0.0s
=> => unpacking to docker.io/library/california-housing-app:latest              212.9s

```

View build details: [docker-desktop://dashboard/build/desktop-linux/desktop-linux/qwusdbntwkliidc8lh2fnxprzeq](#)
(venv) ..\Machine Learning\House-Price-Prediction-MLOPS-End2End-Pipeline> ||

train_test_split
Aa ab .* No results
↑ ↓ ≡ ×

```

1 import requests
2
3 url = "http://localhost:8000/predict"
4
5 payload = {
6     "MedInc": 8.3252,
7     "HouseAge": 41.0,
8     "AveRooms": 6.9841,
9     "AveBedrms": 1.0238,
10    "Population": 322.0,
11    "AveOccup": 2.5556,
12    "Latitude": 37.88,
13    "Longitude": -122.23
14 }
15
16 try:
17     response = requests.post(url, json=payload)
18     response.raise_for_status()
19     print("✅ API responded successfully!")
20     print("📥 Input:", payload)
21     print("📤 Response:", response.json())
22 except requests.exceptions.RequestException as e:
23     print("❌ Error communicating with the API:")
24     print(e)
25

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```

(venv) ..\Machine Learning\House-Price-Prediction-MLOPS-End2End-Pipeline> py tester.py
✅ API responded successfully!
📥 Input: {'MedInc': 8.3252, 'HouseAge': 41.0, 'AveRooms': 6.9841, 'AveBedrms': 1.0238, 'Population': 322.0, 'AveOccup': 2.5556, 'Latitude': 37.88, 'Longitude': -122.23}
📤 Response: {'predicted_price': 4.149}
(venv) ..\Machine Learning\House-Price-Prediction-MLOPS-End2End-Pipeline>

```

```
(venv) ..\Machine Learning\House-Price-Prediction-MLOPS-End2End-Pipeline> uvicorn "California-Housing-Price-Prediction.api.app:app" --host localhost --port 8000
2025-08-03 16:51:30 | INFO | Model and scaler loaded successfully
INFO: Started server process [8500]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: Uvicorn running on http://localhost:8000 (Press CTRL+C to quit)
2025-08-03 16:51:39 | INFO | Started 'read_form'
2025-08-03 16:51:39 | INFO | Completed 'read_form' in 0.015s
INFO: ::1:60833 - "GET / HTTP/1.1" 200 OK
2025-08-03 16:51:53 | INFO | Started 'predict_from_form'
C:\Users\ahc38\OneDrive\Desktop\venv\venv\Lib\site-packages\sklearn\utils\validation.py:2749: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
2025-08-03 16:51:53 | INFO | Form input values:
{
  "MedInc": 2.0,
  "HouseAge": 2.0,
  "AveRooms": 1.0,
  "AveBedrms": 2.0,
  "Population": 2.0,
  "AveOccup": 1.0,
  "Latitude": 12.0,
  "Longitude": 21.0
}
2025-08-03 16:51:53 | INFO | Prediction Output:
{
  "predicted_price": 2.97
}
2025-08-03 16:51:53 | INFO | Completed 'predict_from_form' in 0.021s
2025-08-03 16:51:53 | INFO | Prediction endpoint completed in 0.021s
INFO: ::1:60838 - "POST /form HTTP/1.1" 200 OK
```

California House Price Predictor

Median Income

2

House Age

2

Average Rooms

1

Average Bedrooms

2

Population

2

Average Occupancy

1

Latitude

12

Longitude

21

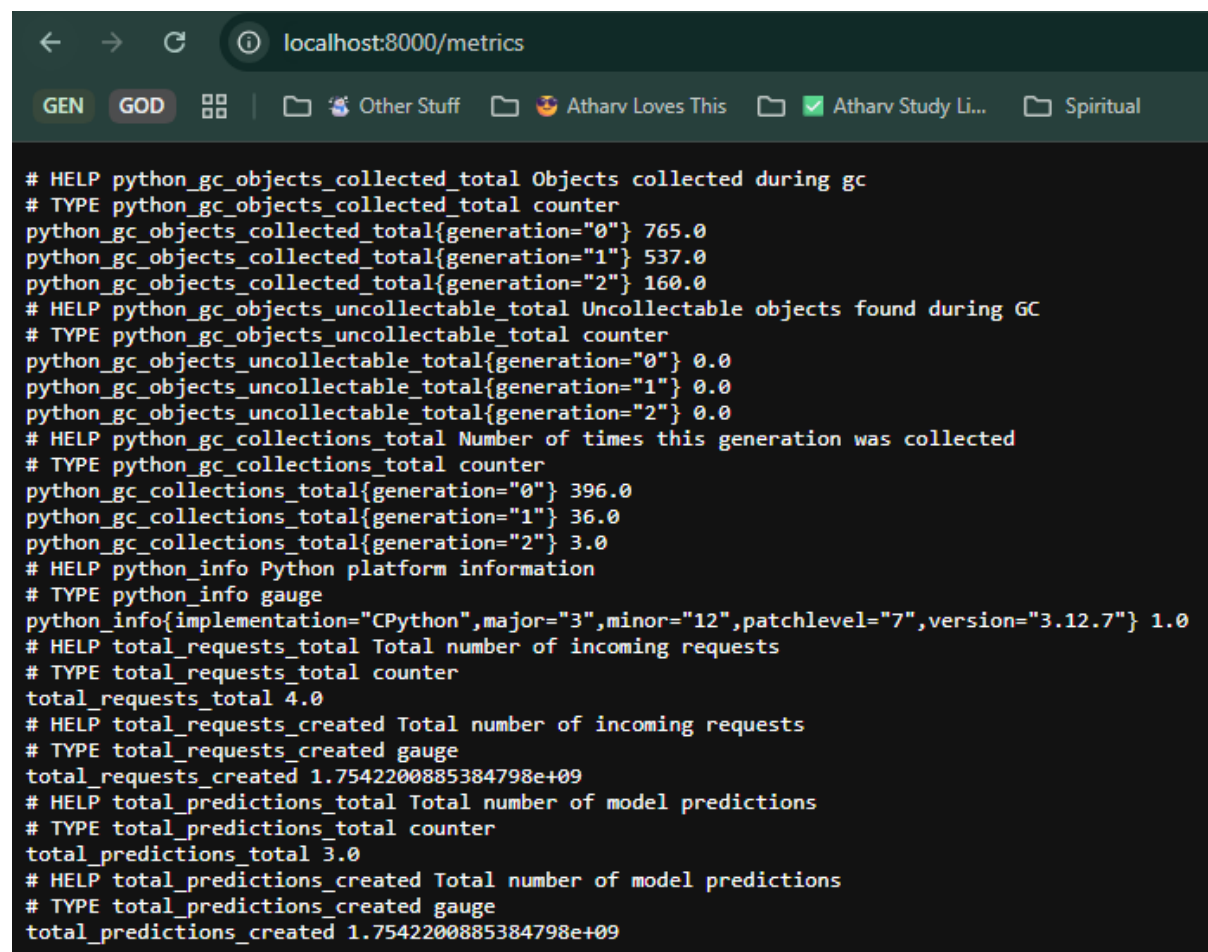
Predict Price

Predicted House Price

\$2.97k

Prediction made at: 2025-08-03 16:53:05

Metrics and Prometheus client Usage :-



The screenshot shows a web browser window with the address bar displaying 'localhost:8000/metrics'. The browser's tab bar includes several tabs: 'GEN', 'GOD', and a folder icon. The main content area displays a list of Prometheus metrics in a dark-themed font. The metrics are organized into groups, each starting with a '# HELP' line followed by a description, a '# TYPE' line indicating the metric type, and then the metric name and its current value. The metrics include: 'python_gc_objects_collected_total' (counter) with values for generations 0, 1, and 2; 'python_gc_objects_uncollectable_total' (counter) with values for generations 0, 1, and 2; 'python_gc_collections_total' (counter) with values for generations 0, 1, and 2; 'python_info' (gauge) showing Python platform information; 'total_requests_total' (counter) with a value of 4.0; 'total_requests_created' (gauge) with a value of 1.7542200885384798e+09; 'total_predictions_total' (counter) with a value of 3.0; and 'total_predictions_created' (gauge) with a value of 1.7542200885384798e+09.

```
# HELP python_gc_objects_collected_total Objects collected during gc
# TYPE python_gc_objects_collected_total counter
python_gc_objects_collected_total{generation="0"} 765.0
python_gc_objects_collected_total{generation="1"} 537.0
python_gc_objects_collected_total{generation="2"} 160.0
# HELP python_gc_objects_uncollectable_total Uncollectable objects found during GC
# TYPE python_gc_objects_uncollectable_total counter
python_gc_objects_uncollectable_total{generation="0"} 0.0
python_gc_objects_uncollectable_total{generation="1"} 0.0
python_gc_objects_uncollectable_total{generation="2"} 0.0
# HELP python_gc_collections_total Number of times this generation was collected
# TYPE python_gc_collections_total counter
python_gc_collections_total{generation="0"} 396.0
python_gc_collections_total{generation="1"} 36.0
python_gc_collections_total{generation="2"} 3.0
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="12",patchlevel="7",version="3.12.7"} 1.0
# HELP total_requests_total Total number of incoming requests
# TYPE total_requests_total counter
total_requests_total 4.0
# HELP total_requests_created Total number of incoming requests
# TYPE total_requests_created gauge
total_requests_created 1.7542200885384798e+09
# HELP total_predictions_total Total number of model predictions
# TYPE total_predictions_total counter
total_predictions_total 3.0
# HELP total_predictions_created Total number of model predictions
# TYPE total_predictions_created gauge
total_predictions_created 1.7542200885384798e+09
```